

## **Exercise: 7**

### **Working with Collections in Cassandra using DataStax Astra**

#### **Aim:**

The aim is to demonstrate how to manage and manipulate various types of collections (sets, lists, maps) in Cassandra through creating, updating, and querying data within a "users" table.

#### **Procedure:**

##### **Initialize Cassandra Table:**

Create a table named users in Cassandra, incorporating diverse data types for columns: an int for user\_id (Primary Key), text for fname and lname, and collections such as a set<text> for emails, a list<text> for top\_places, and a map<timestamp, text> for todo.

##### **Populate Table with Data:**

Sets: Insert a record for a user including a set of email addresses.

Lists: After altering the table to include a list<text> column, update the record to include a list of top places.

Maps: Alter the table to include a map<timestamp, text> column and update the user's record to map tasks to specific timestamps.

##### **Modify Collection Data:**

**Sets:** Demonstrate adding to and removing from a set of emails, including clearing the entire set.

**Lists:** Show how to append to a list, modify a specific element within it, remove an element, and subtract occurrences of specific elements.

**Maps:** Illustrate adding to and removing from a map, including updating specific key-value pairs and clearing a key-value pair.

##### **Cleanup:**

Optionally, demonstrate the deletion of specific columns or the entire row for a user to showcase how to manage and clean up data within Cassandra, ensuring the database remains organized and efficient.

##### **Query and Retrieve Data:**

Execute queries to retrieve updated collections (set, list, map) from the users table for specific user\_id values to demonstrate the changes made during the modification steps.

# Queries :

## Collections – SETS:

A set is an unordered set of values. It is possible to solve the problem of multi-valued fields, like emails.

**1. create a table named users in Cassandra with a column named emails as a set of text Values.**

```
CREATE TABLE users (  
    user_id int PRIMARY KEY,  
    fname text,  
    lname text,  
    emails set<text>  
);
```

```
token@cqlsh> describe keyspaces;  
  
datastax_sla    default_keyspace    system_traces  
system_auth     data_endpoint_auth  system_views  
system_schema  system              system_virtual_schema  
  
token@cqlsh> use default_keyspace;  
token@cqlsh:default_keyspace> CREATE TABLE users (  
    ... user_id int PRIMARY KEY,  
    ... fname text,  
    ... lname text,  
    ... emails set<text>  
    ... );
```

**2. insert data into the emails set for a user with user\_id 1234.**

```
INSERT INTO users (user_id, fname, lname, emails)  
VALUES(1234, 'Frodo', 'Baggins', {'f@baggins.com', 'baggins@gmail.com'});
```

```
token@cqlsh:default_keyspace> INSERT INTO users (user_id, fname, lname, emails)  
    ... VALUES(1234, 'Frodo', 'Baggins', {'f@baggins.com', 'baggins@gmail.com'});  
token@cqlsh:default_keyspace> select*from users;  
  
user_id | emails | fname | lname  
-----+-----+-----+-----  
1234 | {'baggins@gmail.com', 'f@baggins.com'} | Frodo | Baggins  
  
(1 rows)
```

**3. Retrieve the user\_id and emails for the user with user\_id 1234.**

```
SELECT user_id, emails FROM users WHERE user_id = 1234;
```

```
token@cqlsh:default_keyspace> SELECT user_id, emails FROM users WHERE user_id = 1234;
```

user_id	emails
1234	{'baggins@gmail.com', 'f@baggins.com'}

#### 4. Add the email address 'fb@friendsofmordor.org' to the set of emails for the user with user\_id 1234

UPDATE users SET emails = emails + {'fb@friendsofmordor.org'} WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> UPDATE users SET emails = emails + {'fb@friendsofmordor.org'} WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname
1234	{'baggins@gmail.com', 'f@baggins.com', 'fb@friendsofmordor.org'}	Frodo	Baggins

#### 5. Retrieve the user\_id and updated emails for the user with user\_id 1234 after adding the new email.

SELECT user\_id, emails FROM users WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> SELECT user_id, emails FROM users WHERE user_id = 1234;
```

user_id	emails
1234	{'baggins@gmail.com', 'f@baggins.com', 'fb@friendsofmordor.org'}

#### 6.Remove the email address 'fb@friendsofmordor.org' from the set of emails for the user with user\_id 1234

UPDATE users SET emails = emails - {'fb@friendsofmordor.org'} WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> UPDATE users SET emails = emails - {'fb@friendsofmordor.org'} WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname
1234	{'baggins@gmail.com', 'f@baggins.com'}	Frodo	Baggins

#### 7.Remove all email addresses from the set for the user with user\_id 1234.

UPDATE users SET emails = {} WHERE user\_id = 1234;

```
(1 rows)
token@cqlsh:default_keyspace> UPDATE users SET emails = {} WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname
1234	null	Frodo	Baggins

(1 rows)

## Collections – LISTS:

A list is a sorted collection of non-unique values where elements are ordered by their positions in the list.

### 1. Alter the users table to include a column named top\_places as a list of text values

ALTER TABLE users ADD top\_places list<text>;

```
(1 rows)
token@cqlsh:default_keyspace> ALTER TABLE users ADD top_places list<text>;
token@cqlsh:default_keyspace> select*from users;

user_id | emails | fname | lname | top_places
-----+-----+-----+-----+-----
1234 | null | Frodo | Baggins | null
```

### 2. Insert a list of the best places visited into the top\_places column for the user with user\_id 1234.

UPDATE users SET top\_places = [ 'rivendell', 'rohan' ] WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> UPDATE users SET top_places = [ 'rivendell', 'rohan' ] WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;

user_id | emails | fname | lname | top_places
-----+-----+-----+-----+-----
1234 | null | Frodo | Baggins | ['rivendell', 'rohan']

(1 rows)
```

### 3. Add the place 'mordor' to the end of the list of top\_places for the user with user\_id 1234

UPDATE users SET top\_places = top\_places + [ 'mordor' ] WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> UPDATE users SET top_places = top_places + [ 'mordor' ] WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;

user_id | emails | fname | lname | top_places
-----+-----+-----+-----+-----
1234 | null | Frodo | Baggins | ['rivendell', 'rohan', 'mordor']

(1 rows)
```

### 4. Update the second place in the list top\_places to 'riddermark' for the user with user\_id

1234.

UPDATE users SET top\_places[1] = 'riddermark' WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> ALTER TABLE users ADD top_places_map map<int, text>;
token@cqlsh:default_keyspace> select*from users;

user_id | emails | fname | lname | top_places | top_places_map
-----+-----+-----+-----+-----+-----
1234 | null | Frodo | Baggins | ['rivendell', 'rohan', 'mordor'] | null

(1 rows)
token@cqlsh:default_keyspace> UPDATE users SET top_places_map[1] = 'riddermark' WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;

user_id | emails | fname | lname | top_places | top_places_map
-----+-----+-----+-----+-----+-----
1234 | null | Frodo | Baggins | ['rivendell', 'rohan', 'mordor'] | {1: 'riddermark'}

(1 rows)
token@cqlsh:default_keyspace> █
```

## 5.Remove the first place from the top\_places list for the user with user\_id 1234

DELETE top\_places[1] FROM users WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> DELETE top_places_map[1] FROM users WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname	top_places	top_places_map
1234	null	Frodo	Baggins	['rivendell', 'rohan', 'mordor']	null

(1 rows)

## 6.Retrieve the user\_id and top\_places list for the user with user\_id 1234.

SELECT user\_id, top\_places FROM users WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> SELECT user_id, top_places FROM users WHERE user_id = 1234;
```

user_id	top_places
1234	['rivendell', 'rohan', 'mordor']

(1 rows)

## Collections – MAP:

A map allows us to associate two elements, in the form of a key / value. It can be used, for example, to save the schedules of different events in a user profile. Using Cassandra, each element in a Map is stored as a column that we can modify, replace and query.

### 1.Alter the users table to include a column named todo as a map with timestamp keys and text values

ALTER TABLE users ADD todo map<timestamp, text>;

```
token@cqlsh:default_keyspace> ALTER TABLE users ADD todo map<timestamp, text>;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname	todo	top_places	top_places_map
1234	null	Frodo	Baggins	null	['rivendell', 'rohan', 'mordor']	null

(1 rows)

### 2. Insert tasks into the todo map for the user with user\_id 1234.

UPDATE users SET todo = { '2012-9-24' : 'enter mordor', '2012-10-2 12:00' : 'throw ring into mount doom' } WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> UPDATE users SET todo = { '2012-09-24' : 'enter mordor', '2012-10-02T12:00:00.000Z' : 'throw ring into mount doom' } WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname	todo	top_places	top_places_map
1234	null	Frodo	Baggins	{'2012-09-24 00:00:00.000000+0000': 'enter mordor', '2012-10-02 12:00:00.000000+0000': 'throw ring into mount doom'}	['rivendell', 'rohan', 'mordor']	null

(1 rows)

### 3. Add a task with the timestamp '2012-10-2 12:00' to the todo map for the user with user\_id 1234

UPDATE users SET todo['2012-10-2 12:00'] = 'throw my precious into mount doom'  
WHERE  
user\_id = 1234;

```
token@cqlsh:default_keyspace> UPDATE users SET todo['2012-10-2 12:00'] = 'throw my precious into mount doom' WHERE
... user_id = 1234;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname	todo	top_places	top_places_map
1234	null	Frodo	Baggins	{'2012-09-24 00:00:00.000000+0000': 'enter mordor', '2012-10-02 12:00:00.000000+0000': 'throw my precious into mount doom'}	['rivendell', 'rohan', 'mordor']	null

(1 rows)

### 4. Insert tasks into the todo map for the user with user\_id 1234 using the INSERT statement.

INSERT INTO users (user\_id, todo) VALUES ( 1234, { '2013-9-22 12:01' : 'birthday wishes to  
Bilbo', '2013-10-1 18:00' : 'Check into Inn of Prancing Pony' });

```
... 0100 , '2013-10-1 18:00' : 'Check into Inn of Prancing Pony' }));
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname	todo	top_places	top_places_map
1234	null	Frodo	Baggins	{'2013-09-22 12:01:00.000000+0000': 'birthday wishes to Bilbo', '2013-10-01 18:00:00.000000+0000': 'Check into Inn of Prancing Pony'}	['rivendell', 'rohan', 'mordor']	null

(1 rows)

### 5. Remove the task with the timestamp '2012-9-24' from the todo map for the user with user\_id 1234

DELETE todo['2012-9-24'] FROM users WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> DELETE todo['2013-09-22 12:01:00.000000+0000'] FROM users WHERE user_id = 1234;
token@cqlsh:default_keyspace> select*from users;
```

user_id	emails	fname	lname	todo	top_places	top_places_map
1234	null	Frodo	Baggins	{'2013-10-01 18:00:00.000000+0000': 'Check into Inn of Prancing Pony'}	['rivendell', 'rohan', 'mordor']	null

(1 rows)

## 6.Retrieve the user\_id and todo map for the user with user\_id 1234.

SELECT user\_id, todo FROM users WHERE user\_id = 1234;

```
token@cqlsh:default_keyspace> SELECT user_id, todo FROM users WHERE user_id = 1234;

user_id | todo
-----+-----
1234    | {'2013-10-01 18:00:00.000000+0000': 'Check into Inn of Prancing Pony'}

(1 rows)
token@cqlsh:default_keyspace> █
```

## **RESULT:**

The algorithm demonstrates the management of set, list, and map collections in Cassandra, from initialization and data population to modification and retrieval, concluding with optional data cleanup.